

Lung Cancer Survival Prediction Model

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
from xgboost import XGBClassifier
```

```
In [21]: # Load the processed dataset
df = pd.read_csv("processed_dataset_med.csv")
```

```
In [22]: # Function to preprocess non-numeric columns
def preprocess_data(df):
    for col in df.select_dtypes(include=["object"]).columns:
        try:
            # Convert date columns to numerical timestamps
            df[col] = pd.to_datetime(df[col]).astype(int) / 10**9
        except Exception:
            # Convert categorical text columns to numerical categories
            df[col] = df[col].astype("category").cat.codes
    return df

# Apply preprocessing
df = preprocess_data(df)
```

```
In [23]: # Define features (X) and target (y)
X = df.drop(columns=["survived"]) # Assuming 'survived' is the target column
y = df["survived"]
```

```
In [24]: # Split data into training (80%) and testing (20%) sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta
```

```
In [25]: # Scaling the features for better performance
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [26]: # Train Logistic Regression model
log_model = LogisticRegression()
log_model.fit(X_train, y_train)
```

```
Out[26]: ▾ LogisticRegression ⓘ ?
LogisticRegression()
```

```
In [27]: # Predict and evaluate Logistic Regression
y_pred_log = log_model.predict(X_test)
print("Logistic Regression Model:")
```

```
print(f"Accuracy: {accuracy_score(y_test, y_pred_log):.2f}")
print("Classification Report:\n", classification_report(y_test, y_pred_log))
```

Logistic Regression Model:

Accuracy: 0.78

Classification Report:

	precision	recall	f1-score	support
0	0.78	1.00	0.88	138694
1	0.00	0.00	0.00	39306
accuracy			0.78	178000
macro avg	0.39	0.50	0.44	178000
weighted avg	0.61	0.78	0.68	178000

c:\Users\amanv\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

c:\Users\amanv\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

c:\Users\amanv\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

In [28]: *# Train XGBoost model*

```
xgb_model = XGBClassifier(n_estimators=100, learning_rate=0.1, max_depth=6, random_state=42)
xgb_model.fit(X_train, y_train)
```

Out[28]:

XGBClassifier

XGBClassifier(base_score=None, booster=None, callbacks=None, colsample_bylevel=None, colsample_bynode=None, colsample_bytree=None, device=None, early_stopping_rounds=None, enable_categorical=False, eval_metric=None, feature_types=None, gamma=None, grow_policy=None, importance_type=None, interaction_constraints=None, learning_rate=0.1, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=None, min_child_weight=None, missing=None, multi_output_type='raw', num_parallel_tree=None, n_jobs=None, n_estimators=None, num_features=None, num_parallel_features=None, num_threads=None, objective=None, random_state=None, raw_score=False, scale_pos_weight=None, seed=None, silent=None, subsample=None, tree_method=None, verbosity=None, watchlist=None, weight_posterior_scaling=None)

In [29]: *# Predict and evaluate XGBoost*

```
y_pred_xgb = xgb_model.predict(X_test)
print("\nXGBoost Model:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_xgb):.2f}")
print("Classification Report:\n", classification_report(y_test, y_pred_xgb))
```

XGBoost Model:

Accuracy: 0.78

Classification Report:

	precision	recall	f1-score	support
0	0.78	1.00	0.88	138694
1	0.00	0.00	0.00	39306
accuracy			0.78	178000
macro avg	0.39	0.50	0.44	178000
weighted avg	0.61	0.78	0.68	178000

```
c:\Users\amanv\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
c:\Users\amanv\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
c:\Users\amanv\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```