

Mobile Phone Pricing Model

```
In [34]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv("dataset.csv") # Load your dataset
```

```
In [31]: ## Summary of the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   battery_power         2000 non-null   int64
1   blue                  2000 non-null   int64
2   clock_speed           2000 non-null   float64
3   dual_sim              2000 non-null   int64
4   fc                    2000 non-null   int64
5   four_g                2000 non-null   int64
6   int_memory            2000 non-null   int64
7   m_dep                 2000 non-null   float64
8   mobile_wt             2000 non-null   int64
9   n_cores               2000 non-null   int64
10  pc                     2000 non-null   int64
11  px_height              2000 non-null   int64
12  px_width              2000 non-null   int64
13  ram                    2000 non-null   int64
14  sc_h                   2000 non-null   int64
15  sc_w                   2000 non-null   int64
16  talk_time              2000 non-null   int64
17  three_g                2000 non-null   int64
18  touch_screen           2000 non-null   int64
19  wifi                   2000 non-null   int64
20  price_range            2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.3 KB
```

```
In [ ]:
```

```
In [14]: df.describe()
```

Out[14]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_r
count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	1238.518500	0.4950	1.522250	0.509500	4.309500	0.521500	32.000000
std	439.418206	0.5001	0.816004	0.500035	4.341444	0.499662	18.000000
min	501.000000	0.0000	0.500000	0.000000	0.000000	0.000000	2.000000
25%	851.750000	0.0000	0.700000	0.000000	1.000000	0.000000	16.000000
50%	1226.000000	0.0000	1.500000	1.000000	3.000000	1.000000	32.000000
75%	1615.250000	1.0000	2.200000	1.000000	7.000000	1.000000	48.000000
max	1998.000000	1.0000	3.000000	1.000000	19.000000	1.000000	64.000000

8 rows × 21 columns

In [15]: `df.columns`

Out[15]: Index(['battery_power', 'blue', 'clock_speed', 'dual_sim', 'fc', 'four_g',
'int_memory', 'm_dep', 'mobile_wt', 'n_cores', 'pc', 'px_height',
'px_width', 'ram', 'sc_h', 'sc_w', 'talk_time', 'three_g',
'touch_screen', 'wifi', 'price_range'],
dtype='object')

In [16]: `df.shape`

Out[16]: (2000, 21)

In [17]: `df['price_range'].unique()`

Out[17]: array([1, 2, 3, 0])


In [18]: *## To check if we have missing value in the dataset*
`df.isnull().sum()`

```
Out[18]: battery_power    0
         blue             0
         clock_speed     0
         dual_sim        0
         fc              0
         four_g          0
         int_memory      0
         m_dep           0
         mobile_wt       0
         n_cores         0
         pc              0
         px_height       0
         px_width        0
         ram             0
         sc_h            0
         sc_w            0
         talk_time       0
         three_g         0
         touch_screen    0
         wifi            0
         price_range     0
         dtype: int64
```

```
In [19]: ## To check the duplicate records
         df[df.duplicated()]           ## Output : 0 rows x 21 columns
```

```
Out[19]:  battery_power  blue  clock_speed  dual_sim  fc  four_g  int_memory  m_dep  mobile_wt
```

0 rows × 21 columns




```
In [20]: df.corr()
```

Out[20]:

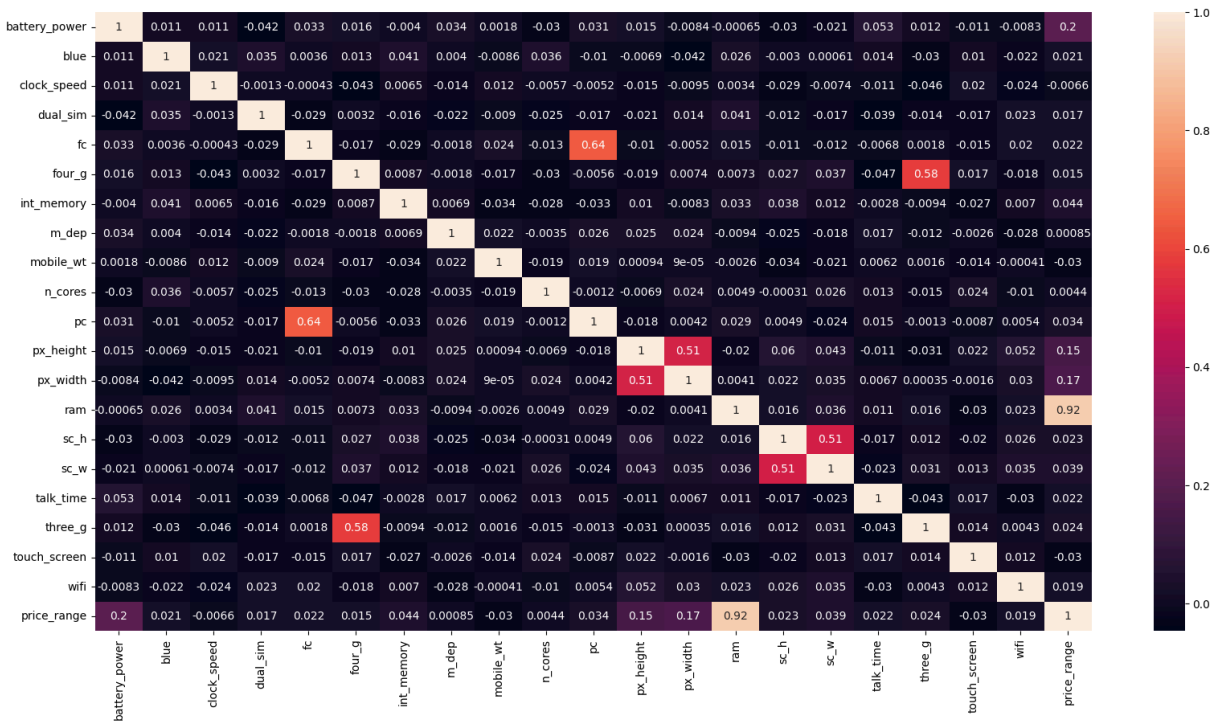
	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_
battery_power	1.000000	0.011252	0.011482	-0.041847	0.033334	0.015665	-
blue	0.011252	1.000000	0.021419	0.035198	0.003593	0.013443	-
clock_speed	0.011482	0.021419	1.000000	-0.001315	-0.000434	-0.043073	-
dual_sim	-0.041847	0.035198	-0.001315	1.000000	-0.029123	0.003187	-
fc	0.033334	0.003593	-0.000434	-0.029123	1.000000	-0.016560	-
four_g	0.015665	0.013443	-0.043073	0.003187	-0.016560	1.000000	-
int_memory	-0.004004	0.041177	0.006545	-0.015679	-0.029133	0.008690	-
m_dep	0.034085	0.004049	-0.014364	-0.022142	-0.001791	-0.001823	-
mobile_wt	0.001844	-0.008605	0.012350	-0.008979	0.023618	-0.016537	-
n_cores	-0.029727	0.036161	-0.005724	-0.024658	-0.013356	-0.029706	-
pc	0.031441	-0.009952	-0.005245	-0.017143	0.644595	-0.005598	-
px_height	0.014901	-0.006872	-0.014523	-0.020875	-0.009990	-0.019236	-
px_width	-0.008402	-0.041533	-0.009476	0.014291	-0.005176	0.007448	-
ram	-0.000653	0.026351	0.003443	0.041072	0.015099	0.007313	-
sc_h	-0.029959	-0.002952	-0.029078	-0.011949	-0.011014	0.027166	-
sc_w	-0.021421	0.000613	-0.007378	-0.016666	-0.012373	0.037005	-
talk_time	0.052510	0.013934	-0.011432	-0.039404	-0.006829	-0.046628	-
three_g	0.011522	-0.030236	-0.046433	-0.014008	0.001793	0.584246	-
touch_screen	-0.010516	0.010061	0.019756	-0.017117	-0.014828	0.016758	-
wifi	-0.008343	-0.021863	-0.024471	0.022740	0.020085	-0.017620	-
price_range	0.200723	0.020573	-0.006606	0.017444	0.021998	0.014772	-

21 rows × 21 columns



In [35]:

```
## Visualising the correlation in the dataset
plt.figure(figsize=(20,10))
sns.heatmap(df.corr(), annot=True)
plt.savefig("correlation_heatmap.png", dpi=300, bbox_inches='tight')
```



```
In [ ]: import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Separate features and target variable
X = df.drop(columns=["price_range"])
y = df["price_range"]

# Split dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st

# Initialize and train the XGBoost model
model = xgb.XGBClassifier(objective="multi:softmax", num_class=4, eval_metric="mlog
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.4f}")
```

```
c:\Users\amanv\AppData\Local\Programs\Python\Python312\Lib\site-packages\xgboost\cor
e.py:158: UserWarning: [18:49:13] WARNING: C:\buildkite-agent\builds\buildkite-windo
ws-cpu-autoscaling-group-i-08cbc0333d8d4aae1-1\xgboost\xgboost-ci-windows\src\leane
r.cc:740:
```

```
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(msg, UserWarning)
```

```
Model Accuracy: 0.9240
```

```
In [23]: # Function to predict price range for new input
def predict_price(new_data):
```

```

new_df = pd.DataFrame([new_data], columns=X.columns)
prediction = model.predict(new_df)[0]
price_categories = {0: "Low Cost", 1: "Medium Cost", 2: "High Cost", 3: "Very H
return price_categories[prediction]

# Example usage
new_mobile = {"battery_power": 1500, "blue": 1, "clock_speed": 2.0, "dual_sim": 1,
              "int_memory": 32, "m_dep": 0.5, "mobile_wt": 140, "n_cores": 4, "pc":
              "px_width": 1200, "ram": 3000, "sc_h": 14, "sc_w": 7, "talk_time": 12
              "touch_screen": 1, "wifi": 1}

predicted_category = predict_price(new_mobile)
print(f"Predicted Price Category: {predicted_category}")

```

Predicted Price Category: Very High Cost

```

In [25]: new_mobile1={
          "battery_power": 4500,"blue": 1,"clock_speed": 3.0,"dual_sim": 1,"fc": 32,"fou
          "int_memory": 32,"m_dep": 0.5,"mobile_wt": 145,"n_cores": 4,"pc": 10,"px_height
          "px_width": 1080,"ram": 12,"sc_h": 12,"sc_w": 6,"talk_time": 14,"three_g": 1,
          "touch_screen": 1,"wifi": 1
        }
predicted_category = predict_price(new_mobile1)
print(f"Predicted Price Category: {predicted_category}")

```

Predicted Price Category: Low Cost

```

In [26]: new_mobile_low = {
          "battery_power": 800, "blue": 0, "clock_speed": 1.0, "dual_sim": 0, "fc": 2, "f
          "int_memory": 8, "m_dep": 0.3, "mobile_wt": 180, "n_cores": 2, "pc": 5,
          "px_area": 500 * 800, # Using px_area instead of px_height & px_width
          "ram": 512, "sc_h": 10, "sc_w": 4, "talk_time": 7, "three_g": 0,
          "touch_screen": 0, "wifi": 0
        }

predicted_category = predict_price(new_mobile_low)
print(f"Predicted Price Category: {predicted_category}")

```

Predicted Price Category: Low Cost


```

In [37]: from sklearn.metrics import confusion_matrix

# Generate confusion matrix
cm = confusion_matrix(y_test, y_pred)

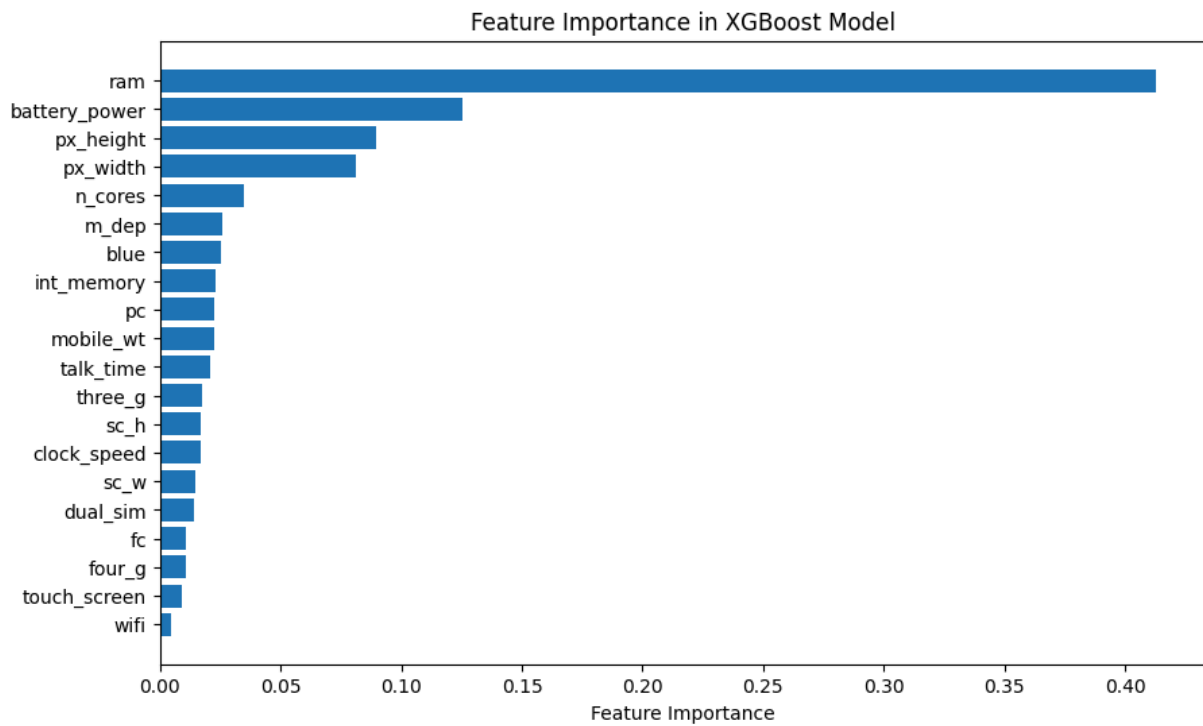
# Plot heatmap
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Low", "Medium", "H
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

```



```
In [28]: importances = model.feature_importances_
feature_names = X.columns

# Sort and plot
sorted_idx = importances.argsort()
plt.figure(figsize=(10,6))
plt.barh([feature_names[i] for i in sorted_idx], importances[sorted_idx])
plt.xlabel("Feature Importance")
plt.title("Feature Importance in XGBoost Model")
plt.show()
plt.savefig("imp_features.png", dpi=200, bbox_inches='tight')
```



<Figure size 640x480 with 0 Axes>

```
In [38]: from sklearn.metrics import classification_report

# Generate classification report
report = classification_report(y_test, y_pred)

# Print report
print(report)

# Save classification report to a text file
with open("classification_report.txt", "w") as f:
    f.write(report)

print("Classification report saved as 'classification_report.txt'")
```

	precision	recall	f1-score	support
0	0.97	0.98	0.98	125
1	0.93	0.90	0.91	125
2	0.85	0.90	0.88	125
3	0.95	0.91	0.93	125
accuracy			0.92	500
macro avg	0.92	0.92	0.92	500
weighted avg	0.92	0.92	0.92	500

Classification report saved as 'classification_report.txt'